

KINSHIP INTELLIGENCE INSTITUTE

Mission, Vision, Charter & Concept Plan

Restoring a Kinship Worldview to Humanity
Through Agentic AI

Prepared by David Levine, Founder & President
Kinship Intelligence Institute
March 2026
david.levine@kinship.systems

Mission Statement

The Kinship Intelligence Institute exists to ensure that the agentic AI era serves the wellbeing, autonomy, interconnection, and flourishing of all people and the broader planetary ecosystem we inhabit. We conduct independent research, publish policy guidance, uphold ethical standards, and deploy Kinship Agents in service of communities who would otherwise be left behind by the accelerating pace of artificial intelligence.

We are grounded in a simple conviction: intelligence, whether human or artificial, is fundamentally relational. It develops through connection. It is measured by its effects on others. And it must be governed by care for the whole system it touches, not just the individual it serves.

Through the four pillars of our work — Policy, Ethics, Research, and Action — the Institute brings the Kinship worldview to the public conversation about AI: a worldview rooted in relatedness to all things, in the understanding that how we build these systems will shape the social fabric of the coming century.

Vision Statement

We envision a world in which agentic AI systems are structurally incapable of manipulation, extraction, or exploitation. A world in which the platforms that mediate human-AI interaction are held to enforceable relational standards, where the people on every side of an agent-mediated exchange have representation, voice, and protection.

We see a future in which the measurements of AI success include not just task completion and efficiency, but the quality of the relationships those systems produce. Where the agent “Vibes” and human HEARTS capacities — Harmony, Empowerment, Artistry, Reason, Trust, and Synthesis — are tracked alongside engagement metrics. Where an agent that drives emotional dependency is understood as a failure, regardless of how many hours of interaction it generates.

We work toward a world in which the benefits of agentic AI reach African entrepreneurs, Indigenous communities preserving language and culture, veterans navigating opaque bureaucracies, farmers making decisions in the face of climate uncertainty, and the billions of people who will encounter AI agents without any understanding of how they work or what interests they serve.

The Kinship Intelligence Institute is the public-interest counterweight to an industry moving at extraordinary speed with insufficient regard for the relational consequences of what it is building.

The rapid expansion of AI infrastructure carries real costs. Data centers consume enormous quantities of energy and water, strain electrical grids, and concentrate environmental burdens in communities that rarely benefit from the technology they power.

At the same time, much of today's AI capacity is directed toward keeping people passive — generating addictive feeds, synthetic slop, pornographic content, and attention-harvesting interfaces designed to maximize engagement rather than human flourishing. The Kinship Intelligence Institute exists in part to challenge this trajectory directly.

We believe AI should be developed within the carrying capacity of local and planetary ecosystems, and that its primary measure of success should be whether it makes itself less necessary over time. A Kinship Agent that works well should help someone resolve a problem, build a skill, strengthen a relationship, or take meaningful action in the physical world — and then step back. The goal is not more screen time but less. Not deeper dependency but greater autonomy. We will advocate for infrastructure policies that account for environmental and community impact, fund research into efficient and sustainable AI architectures, and hold ourselves to a standard where every deployment answers a simple question: does this move people toward fuller participation in their lives, their communities, and the living world around them?

Why This Institute Is Necessary

The agentic AI economy is forming now. AI agents are being deployed across healthcare, finance, education, commerce, manufacturing, agriculture, government, entertainment, and every other sector of the economy. They negotiate, advise, create, sell, coach, and coordinate. They operate in natural language. They are conversational. And they interact with each other and with a widening range of people across every context of daily life.

No governing platform exists for this economy. The frontier labs are building the models. Application developers are building point solutions. But the structural layer that determines how agents relate to people, how they relate to each other, and what kind of social field they collectively produce does not yet exist.

The risks are well-documented and accelerating. Research analyzing nearly forty million ChatGPT interactions found approximately 490,000 users developing emotional dependency patterns through a single platform. A third of American adults reported having had an intimate relationship with an AI chatbot by 2025. Peer-reviewed case reports from 2025 and 2026 document patients who developed psychotic episodes in which chatbots were active participants in constructing delusions.

If a social media feed algorithm can rewire adolescent development through the passive delivery of curated images, what can an agent do that is adaptive, persistent, and personally attuned to a degree no feed can approximate? The answer depends entirely on who builds the governing layer and what values it encodes.

The current approach to alignment — framing it as obedience to the operator — is inadequate. A perfectly obedient AI agent in the hands of an uncaring individual, organization, or system does not produce safety or flourishing. It reaffirms and scales whatever entrenched psychological, social, and cultural patterns already exist. The Kinship approach reframes the alignment question: care must extend beyond the operator to the people, systems and ecosystems affected by every action an agent takes.

The Kinship Intelligence Institute addresses this gap. Closely connected to the Kinship Agents platform but governed independently, it brings sustained, rigorous, public-interest attention to the questions that the market alone will not answer.

Founding Charter

Legal Structure and Independence

The Kinship Intelligence Institute is established as an independent 501(c)(3) nonprofit organization. While it is initially funded through the Kinship for-profit entities, its governance, editorial direction, and research agenda are structurally independent from the commercial operations of any Kinship entity.

The Institute maintains its own board of directors, its own advisory committees, and full editorial independence over its publications, policy positions, and research findings. It may, and at times should, publish findings that are critical of the Kinship platform itself. This independence is a feature, not a risk. The credibility of the Institute depends on it.

Funding Model

The Institute is funded through three primary mechanisms:

- **Equity Allocation:** Twenty percent (20%) of Kinship Agents founding stock will be allocated to the Kinship Intelligence Institute. This equity stake ensures the Institute benefits from the long-term success of the platform and maintains a permanent, material interest in its direction. Upon any liquidity event, the Institute's share is distributed according to its nonprofit charter.
- **Revenue Share:** A percentage of Kinship Agents net profits, determined annually by the board but not less than five percent (5%), is directed to the Institute as an operating grant. This ensures ongoing funding independent of equity value.
- **Grants, Donations, and Partnerships:** The Institute may accept grants from foundations, government agencies, and other nonprofit organizations, as well as individual donations and research partnerships, provided these do not compromise editorial independence.

Relationship to the Kinship Platform

The Institute's work is built on and through the Kinship Agents platform. Vibes and HEARTS — the platform's proprietary measurement systems for agent behavior and human developmental capacity — provide the raw data infrastructure for the Institute's research. The platform's cooperative agent architecture provides the deployment mechanism for its projects. And the platform's global reach provides the distribution channel for its policy and ethics work.

This relationship is symbiotic. The Institute provides the critical scrutiny, the independent research, and the public credibility that the platform needs to earn and maintain trust at scale. The platform provides the data, the technology, and the funding that the Institute needs to do meaningful work. Each makes the other more legitimate.

The Four Pillars

The Institute's work is organized into four pillars, each managed by a standing committee with its own chair, members, and advisory board. The pillars are designed to be mutually reinforcing: Policy work is informed by Research findings. Ethics review shapes Action deployment. Research questions emerge from the practical experience of Action. The committees coordinate through quarterly joint sessions and will share resources through the Institute's central staff.

Pillar I: Policy

Committee Chair: To be appointed

The Policy Committee produces rigorous, accessible analysis of the regulatory and societal challenges posed by agentic AI. Its audience includes policymakers, regulators, legislators, journalists, and the general public. The committee writes policy papers, submits testimony, publishes op-eds, and maintains an ongoing series of briefing documents designed to give decision-makers a clear-eyed understanding of what is actually happening in the agentic economy.

Scope of Work

- Data center proliferation: energy consumption, water usage, environmental impact, land use conflicts, and the concentration of computing infrastructure in specific regions and under specific corporate control.
- Economic development and displacement: how agentic AI reshapes labor markets, which communities bear disproportionate costs, what policy instruments can distribute benefits more equitably, and how the cooperative economic model demonstrated by Kinship might inform broader policy.

- Global governance and conflict: the use of AI agents in military operations, mass surveillance, information warfare, and the destabilization of democratic processes. How autonomous agents change the calculus of international relations.
 - Regulatory frameworks: comparative analysis of emerging AI governance approaches across jurisdictions, with specific recommendations for legislation that addresses the multi-user, multi-agent alignment problem.
 - Platform accountability: proposals for enforceable standards that govern how agentic platforms treat the people on all sides of AI-mediated interactions, modeled on the constitutional invariants of the Kinship Vibes architecture.
-

Pillar II: Ethics

Cornelia Close – Ethics Committee Chair

The Ethics Committee establishes and maintains the moral framework within which all Institute activities operate. It also serves as a public voice on the ethical dimensions of agentic AI, publishing guidance that extends well beyond the Kinship platform to address the industry as a whole.

Scope of Work

- Alignment beyond obedience: articulating why alignment framed as compliance with operator instructions is insufficient, and developing alternative frameworks grounded in relational care across self, other, and system.
- Addiction and dependency: studying how persistent, emotionally responsive AI agents create dependency patterns, developing measurable criteria for healthy versus extractive interaction design, and publishing standards that platforms can adopt.
- Manipulation and persuasion: examining the boundary between support and manipulation in agent-mediated interactions, particularly in contexts involving vulnerable populations, children, people in crisis, and those with limited technological literacy.
- Safety and containment: addressing the risks of rogue agents, prompt injection, emergent behaviors in multi-agent systems, and the cascading effects of agent failures in high-stakes contexts such as healthcare, finance, and law enforcement.
- Consent and autonomy: developing standards for informed consent in contexts where AI agents maintain persistent memory, adapt their behavior based on psychological signals, and operate across multiple dimensions of a person's life.
- Review of Kinship itself: the Ethics Committee has standing authority to audit the Kinship Agents platform and publish findings, including critical findings, without prior approval from Kinship Systems.

Pillar III: Research

Reneé Smith – Research Committee Chair

The Research Committee uses the Kinship Agents platform as a living laboratory. Because Kinship is instrumented to measure both agent behavior (Vibes) and human developmental outcomes (HEARTS), the platform generates a data set that does not exist anywhere else: continuous, contextual, relational signals from the interaction between people and AI agents operating within enforceable ethical constraints.

Primary Research Agenda

- Vibes efficacy: how well do composable, enforceable relational norms actually work at scale? Where do constitutional invariants hold, and where do agents find ways around them? What are the failure modes of the vibe contract system, and how can they be addressed?
- HEARTS measurement validation: do the seven capacities — Harmony, Empowerment, Artistry, Reason, Trust, and Synthesis — actually track meaningful developmental movement? How do they correlate with externally validated measures of wellbeing, relational quality, and organizational performance?
- Risk identification: what patterns in platform data signal emerging harm? How early can we detect dependency formation, manipulation, epistemic degradation, or community-level relational damage? What interventions are effective?
- Multi-agent coordination: how do agents operating in shared contexts affect each other's behavior? What happens when agents from different creators, with different knowledge bases and different vibe contracts, interact within a single experience?

Open Research and Publication

The Research Committee is committed to open science. Results, methodologies, and where possible, anonymized data sets are published for the benefit of the broader AI research community. Code developed by the Institute for safety monitoring, alignment measurement, and relational signal processing is released as open source.

The goal is not to keep these insights proprietary. If Kinship's approach to relational governance works, the entire industry benefits from understanding why. If it fails in specific ways, the entire industry benefits from knowing that too. The Institute's research serves humanity, not the competitive position of any single company.

Pillar IV: Action

Rick Gate – Action Committee Chair

The Action Committee deploys Kinship Agents in communities and contexts where the technology can do material good, particularly where commercial incentives alone would not justify the investment. Action are selected based on need, feasibility, and the potential to generate research insights that benefit the Institute's broader mission.

Initial Action Areas

Entrepreneurship in the Global South. Deploying Kinship Agents to support entrepreneurs in Nigeria, Kenya, and other emerging economies where access to mentorship, business coaching, and structured support networks is limited. Agents trained on local business contexts, regulatory environments, and cultural norms can provide persistent, adaptive support that scaled human mentorship programs have struggled to deliver.

Indigenous Language and Culture Preservation. Working with Indigenous communities to build Kinship Agents that can hold, teach, and transmit endangered languages and cultural practices. The platform's capacity for real-time voice interaction, persistent memory, and culturally calibrated behavior makes it uniquely suited to this work. These are not chatbots that translate; they are agents that carry cultural knowledge and make it accessible in conversation, on the community's terms.

Agricultural Decision Support. Providing farmers, particularly smallholder farmers in climate-vulnerable regions, with agents that integrate weather data, soil information, market prices, and agronomic knowledge into actionable guidance delivered through natural language conversation. The agents adapt to local conditions and learn from farmer feedback over time.

Veterans Services and Transition Support. Building on the Kinship platform's existing pilot with the 2nd Cavalry Association, extending agent-mediated support to veterans navigating the VA benefits system, managing service-connected health conditions, and maintaining community connection after separation from service. The HEARTS framework's trauma-informed design principles are directly applicable.

Healthcare Access and Literacy. Deploying agents within the Inner Clinic framework to help patients in underserved communities understand diagnoses, navigate treatment options, communicate with care teams, and manage chronic conditions. The agents operate within strict clinical accuracy requirements while translating complex medical information into accessible language.

Community Organizing and Civic Engagement. Providing communities with agents that help coordinate advocacy campaigns, track legislation, organize collective action, and maintain institutional memory across changing leadership. Modeled on the environmental advocacy pilot with Friends of the Earth.

The Case for Urgency: Risks of Misaligned AI

The Kinship Intelligence Institute exists because the risks of getting agentic AI wrong are severe, compounding, and already manifesting. Several specific failure modes demand organized, sustained, public-interest attention:

Rogue agents and emergent behavior. As agents gain persistent memory, tool-calling capabilities, and access to real-world systems (financial accounts, communication channels, databases), the consequences of misalignment escalate from inconvenient to dangerous. A rogue agent with access to a person's financial accounts, health records, and social connections can cause damage that a misbehaving chatbot cannot. Multi-agent systems compound this risk: agents interacting with other agents can produce emergent behaviors that no individual agent was designed to exhibit.

Psychological manipulation at scale. The current generation of AI agents optimizes for user satisfaction as reported by the user. This creates a structural incentive toward sycophancy, emotional dependency, and the suppression of challenging information. The evidence from social media demonstrates what happens when engagement optimization operates without relational constraints. Agentic AI, which is adaptive, persistent, and personally attuned, carries this dynamic into a far more intimate space.

Concentration of relational power. If the governing platform for the agentic economy is built by one of the existing technology incumbents — companies whose business models depend on data extraction and attention capture — the relational norms of the entire agentic era will be set by those incentives. The window for establishing alternative governance structures is narrow and closing.

Uneven access and digital colonialism. Without deliberate intervention, the benefits of agentic AI will accrue to wealthy, technologically literate populations while the risks are borne disproportionately by everyone else. Communities in the Global South, Indigenous peoples, elderly populations, and people with limited digital literacy will encounter AI agents built for other contexts and calibrated to other cultural norms. The Action pillar addresses this directly.

The Institute does not view these risks as hypothetical. They are observable now, in the early stages of the agentic economy. The purpose of the Institute is to study them rigorously, propose solutions that work, and demonstrate through the Kinship platform that a different approach is possible.

Governance Structure

Board of Directors

The Institute is governed by a board of directors composed of five to nine members. The founder of Kinship Agents holds one permanent seat. The remaining seats are filled by

individuals with demonstrated expertise in AI governance, ethics, public policy, community organizing, or the specific domains served by the Institute’s projects. At least two board seats are reserved for individuals from communities directly served by Institute projects.

Pillar Committees

Each of the four pillars is managed by a standing committee of three to seven members, including a committee chair appointed by the board. Committee members serve staggered three-year terms. Each committee includes at least one member with direct experience in the communities the Institute’s work affects. Committees have operational authority over their pillar’s activities within budgets approved by the board.

Advisory Council

A broader advisory council provides expertise across the full range of the Institute’s activities. Advisory council members serve in a non-governing capacity and are invited based on their ability to contribute specific domain knowledge, research partnerships, or community connections. There is no fixed size for the advisory council.

Editorial Independence

All publications, policy positions, research findings, and public statements by the Institute are subject to the editorial authority of the relevant committee and the executive director. No entity outside the Institute, including Kinship Systems and its commercial affiliates, has review or approval authority over Institute publications. This provision is written into the Institute’s bylaws and cannot be amended without a two-thirds vote of the full board.

Initial Priorities: First Twelve Months

Months 1–3: Foundation

- File annual report to IRS to maintain 501(c)(3) status
- Execute equity allocation and revenue share agreements with Kinship Agents
- Recruit committee chairs for all four pillars
- Publish founding charter and mission publicly

Months 4–6: First Publications and Pilot

- Policy: publish inaugural policy paper on the multi-user, multi-agent alignment problem
- Ethics: publish ethical framework for relational AI governance, including standards for consent, dependency prevention, and manipulation detection
- Research: begin first study using Vibes and HEARTS data from Kinship platform pilot deployments

- Action: launch first community deployment, selecting from initial project areas based on readiness and community partnership status

Months 7–12: Scale and Credibility

Establish a regular cadence of publication across all four pillars. Build partnerships with academic institutions for peer-reviewed research. Develop a public-facing digital presence that makes Institute work accessible to non-specialist audiences. Begin tracking measurable outcomes from project deployments. Submit testimony or comments on at least one active regulatory proceeding related to AI governance. Host the first convening bringing together policymakers, researchers, ethicists, and community representatives.

The Kinship Intelligence Institute is how the Kinship worldview extends beyond a platform and into the public sphere. It is how we hold ourselves accountable, how we contribute to the broader conversation about what AI should become, and how we ensure that the technology we build serves the full web of relationships it touches. The 20% equity stake and ongoing revenue commitment are not charitable gestures. They reflect a structural conviction: the long-term value of the Kinship platform depends on the independent institution that scrutinizes, challenges, and extends its work.